

Autonomous Agents 1

Assignment 2

By Group 4: Gieske, Gornishka, Koster, Loor

November 28, 2014

Contents

1	Introduction	3
2	Theory	4
2.1	Temporal difference learning methods	4
2.2	Q-Learning	4
2.3	Sarsa	5
2.4	Monte Carlo Methods	5
2.5	On-policy Monte Carlo Control	5
2.6	Off-policy Monte Carlo Control	6
2.7	Action selection methods	7
3	Implementation	9
4	Analysis	10
4.1	Q-Learning	10
4.1.1	Different discount factors	10
4.1.2	Different learning rates	10
4.1.3	The effect of ϵ	11
4.1.4	Softmax action selection	11
4.1.5	ϵ -greedy vs. softmax	12
4.1.6	Different initialization for Q-learning	13
4.2	Learning types	14
5	Conclusion	17
5.1	Parameter settings	17
6	Future work	18
6.1	Encoded version	18
6.2	Learning rate	18
6.3	Varying grid initialization	18
6.4	Multiple experiments	18
6.5	Exploration-Exploitation optimization	18
7	Files attached	19
8	Sources	19

1 Introduction

This report contains the analysis of an implementation of single agent reinforcement learning. In these algorithms, the agent has no prior knowledge of the transition or reward function. Two different solutions to this problem are presented: in model-based learning, the agent learns the model from experience and plans on the learned model. In model-free learning, the agent learns state-values or Q-values directly. In this paper, model-free learning methods are considered.

Some prominent model-free reinforcement learning methods are on- and off-policy Monte Carlo Control, Q-learning (off-policy) and Sarsa (on-policy), where on-policy methods evaluate the values of a policy π *as it's being followed*, while off-policy methods allow the agent to follow a different policy π' . This paper reports on Q-learning, its on-policy equivalent, Sarsa, as well as on On- and Off-Policy Monte Carlo Control methods. Moreover, the difference in action-selection between ϵ -greedy and softmax are compared.

2 Theory

2.1 Temporal difference learning methods

Temporal difference (TD) learning methods estimate state-values or Q-values by updating them on every time step (as opposed to the episodic updates of Monte Carlo-methods). Temporal difference learning methods are similar to Monte Carlo methods in that they both learn from experience directly, without requiring a model and they are similar to dynamic programming methods in that they both exploit the Bellman-equations and *bootstrap*, i.e. they estimate values based on earlier estimates, without waiting for the final outcome. However, since TD-methods do not require a model, the state- and action-value updates are slightly different than shown in 1 and 2. However, as TD-methods do not require a model, the state- and action value updates are different than shown in 1 and 2. The value update rule is displayed in 3. Updating the Q-values differs per learning method and are shown per section.

$$V^*(s) = \sum_{s'} \mathcal{P}_{ss'}^a \left[\mathcal{R}_{ss'}^a + \gamma \max_a V^*(s') \right] \quad (1)$$

$$Q^*(s, a) = \sum_{s'} \mathcal{P}_{ss'}^a \left[\mathcal{R}_{ss'}^a + \gamma \max_{a'} Q^*(s', a') \right] \quad (2)$$

Thus, in general, an update-rule for a TD-method would be similar to that of TD(0) in equation 3.

$$V(s) \leftarrow V(s) + \alpha[r + \gamma V(s') - V(s)] \quad (3)$$

Both Q-learning and Sarsa are TD-methods, and are described in more detail in sections Q-learning and Sarsa.

2.2 Q-Learning

Q-learning¹ is a temporal difference method that uses the update rule in equation 4. Because it retrieves the Q-value of the state-action pair where $Q(s', a)$ is maximized, it is an *off-policy* method. The algorithm for Q-learning can be found in pseudocode in figure 1.

$$Q(s_t, a_t) \leftarrow Q(s_t, a_t) + \alpha \left[r_{t+1} + \gamma \max_a Q(s_{t+1}, a) - Q(s_t, a_t) \right] \quad (4)$$

```

Initialize Q(s,a) arbitrarily
Repeat (for each episode):
  Initialize s
  Repeat (for each step of episode):
    Choose a from s' using policy derived from Q (e.g.,  $\epsilon$ -greedy)
    Take action a, observe r, s'
     $Q(s,a) \leftarrow Q(s,a) + \alpha[r + \gamma \max_{a'} Q(s', a') - Q(s, a)]$ 
     $s \leftarrow s'$ ;
  until s is terminal

```

Figure 1: The algorithm for one-step Q-learning [2]

¹More specifically, this is *one-step Q-learning*, which is the algorithm evaluated in this paper.

2.3 Sarsa

Sarsa is a temporal difference method that uses the update rule in equation 5. Because it retrieves the Q-value of the state-action pair (s', a') where a' is selected using the policy π that is being evaluated, it is an *on-policy* method. The algorithm for Sarsa can be found in pseudocode in figure 2.

$$Q(s_t, a_t) \leftarrow Q(s_t, a_t) + \alpha [r_{t+1} + \gamma Q(s_{t+1}, a_{t+1}) - Q(s_t, a_t)] \quad (5)$$

```
Initialize Q(s,a) arbitrarily
Repeat (for each episode):
  Initialize s
  Choose a from s' using policy derived from Q (e.g.,  $\epsilon$ -greedy)
  Repeat (for each step of episode):
    Take action a, observe r, s'
    Choose a' from s' using policy derived from Q (e.g.,  $\epsilon$ -greedy)
     $Q(s,a) \leftarrow Q(s,a) + \alpha[r + \gamma Q(s', a') - Q(s, a)]$ 
     $s \leftarrow s'$ 
  until s is terminal
```

Figure 2: The algorithm for Sarsa [2]

2.4 Monte Carlo Methods

As stated before, there are several ways of learning an estimating value functions and discovering optimal policies. Again, Monte Carlo methods do not know (nor need) complete knowledge of the environment. Monte Carlo methods require only experience, such as sample sequences of states, actions, and rewards, from on-line interaction with an environment. Unlike TD-methods, Monte Carlo methods require a model. The model only needs to generate sample transitions, unlike Dynamic Programming methods which need complete probability distributions of all possible transitions. This is a big advantage, since obtaining the whole probability distribution might be a tedious, even infeasible task in cases when sampling is still possible. Therefore, Monte Carlo methods are an important tool for an agent to use when only the model is known.

2.5 On-policy Monte Carlo Control

On-policy Monte Carlo control (ONMC) does not use exploring starts. In order to ensure that each state is visited and each action is selected infinitely often is to keep selecting these actions. This is done by using an ϵ -greedy algorithm, which guarantees exploration. There are two ways of implementing ONMC: first-visit Monte Carlo and every-visit Monte Carlo. In this case, every-visit Monte Carlo is implemented which is reflected in the algorithm below.

```

Initialize, for all  $s \in (S), a \in (A)(s)$ 
   $Q(s,a) \leftarrow$  arbitrary
   $Returns(s,a) \leftarrow$  empty list
   $\pi \leftarrow$  an arbitrary  $\epsilon$ -soft policy
Repeat forever:
  Generate an episode using exploring starts and  $\pi$ 
  For each pair  $s,a$  appearing in the episode:
     $\mathcal{R} \leftarrow$  return following the first occurrence of  $s,a$ 
    Append  $\mathcal{R}$  to  $Returns(s,a)$ 
     $Q(s,a) \leftarrow \text{average}(Returns(s,a))$ 
  For each  $s$  in the episode:
     $a^* \leftarrow \arg \max_a Q(s,a)$ 
    For all  $a \in \mathcal{A}(s)$ :
      
$$\pi(s,a) \leftarrow \begin{cases} 1 - \epsilon + \frac{\epsilon}{|\mathcal{A}(s)|} & \text{if } a = a^* \\ \frac{\epsilon}{|\mathcal{A}(s)|} & \text{if } a \neq a^* \end{cases}$$


```

Figure 3: The algorithm for on-policy Monte Carlo [2]

2.6 Off-policy Monte Carlo Control

Similarly to ONMC, off-policy Monte Carlo control (OFFMC) does not use exploring starts. Thus, an ϵ -greedy algorithm is used again to ensure exploration. The main difference between OFFMC and ONMC is in that ONMC estimates the value of a policy while using it for control (i.e. exploration and exploitation are done together), whereas OFFMC uses two different policies. The first one is called behavior policy and it is only used to generate behavior (i.e. this policy is used in the exploration part). The second one is called an estimation policy and it's the one that is being evaluated and improved (i.e this policy is used in the exploitation part). There are again two ways of implementing OFMCC: first-visit Monte Carlo and every-visit Monte Carlo. In this case, every-visit Monte Carlo is implemented which is reflected in the algorithm in figure 4.

First, an episode is generated and its corresponding state,action pairs are stored. If the greedy-policy would select a different action than the behavioral policy, the time-step t at which this happens is stored, to ensure that later on in the algorithm, we only evaluate the Q-values of state-action pairs that follow the deterministic (greedy) policy. This is done because when the weights w are calculated, the probability of the deterministic action is divided by the probability of the behavioral action. However, the first is either zero (if it is not the greedy action) or one (if it is the greedy action). Thus, storing t will save a lot of computation time.

The, for each visited pair in the episode, the weight is calculated, and the nominator and denominator are updated by adding wR_t (where R_t is the return) and w respectively. The Q-value is then $\frac{\text{nominator}}{\text{denominator}}$. The nominator, denominator and Q-values are stored, and the computed Q-values are used to update the deterministic policy greedily according to $\pi(s) \leftarrow \arg \max_a Q(s,a)$.

```

Initialize, for all  $s \in \mathcal{S}, a \in \mathcal{A}(s)$ :
     $Q(s,a) \leftarrow$  arbitrary
     $N(s,a) \leftarrow 0$ ; Numerator
     $D(s,a) \leftarrow 0$ ; Denominator of  $Q(s,a)$ 
     $\pi \leftarrow$  an arbitrary deterministic policy
Repeat forever:
    Select a policy  $\pi'$  and use it to generate an episode:
         $s_0, a_0, r_1, s_1, r_2, \dots, s_{T-1}, a_{T-1}, r_T, s_T$ 
     $\tau \leftarrow$  latest time at which  $a_t \neq \pi(s_\tau, a_\tau)$ 
    For each pair  $s, a$  appearing in the episode at time  $\tau$  or later:
         $\tau \leftarrow$  the time of first occurrence of  $s, a$  such that  $t \leq \tau$ 
         $w \leftarrow \prod_{k=\tau+1}^{T-1} \frac{1}{\pi'(s_k, a_k)}$ 
         $N(s,a) \leftarrow N(s,a) + wR_t$ 
         $D(s,a) \leftarrow D(s,a) + w$ 
         $Q(s,a) \leftarrow \frac{N(s,a)}{D(s,a)}$ 
    For each  $s \in \mathcal{S}$ :
         $\pi(s) \leftarrow \arg \max_a Q(s,a)$ 

```

Figure 4: The algorithm for off-policy Monte Carlo [2]

2.7 Action selection methods

In order to select which action to choose according to a given policy, a tradeoff between exploration and exploitation must take place. This tradeoff is important when performing reinforcement learning as the rewards must be maximized, but exploration may lead to finding higher rewards. There are several action selection methods which can be used to select actions. The two techniques analyzed in this report are ϵ -greedy and softmax action selection.

In the case of ϵ -greedy, the best action a^* is given a probability of $1 - \epsilon$. All actions a (including a^*) then receive an equal portion of ϵ as a probability. This is formalized in equation 6.

$$\forall a \in \mathcal{A} : p(a) = \begin{cases} 1 - \epsilon + \frac{\epsilon}{|\mathcal{A}|}, & \text{if } a = a^* \\ \frac{\epsilon}{|\mathcal{A}|} & \text{otherwise} \end{cases} \quad (6)$$

So that in $1 - \epsilon$ of the cases, a^* is selected, and in ϵ of the cases a random action is chosen uniformly. Thus, an ϵ -greedy policy mainly exploits the known states to maximize the immediate reward. At probabilistically determined times however, this policy will explore new states. This may lead to undesired behavior as it is possible for the agent to stand beside the goal state when the ϵ -greedy policy turns to explore a new state, which can lead to high negative rewards in particular cases, such as a robot falling off a cliff. Also, ϵ -greedy does not consider the quality of non-optimal actions: an action with a value just below that of a^* receives the same probability as an action with a much lower value. Another note using ϵ -greedy is the intuition that ϵ should decay as the agent has explored more states, as exploration is adding less information over time. There is, however, no clear-cut way to decide how and when to decay ϵ .

Softmax action-selection offers a solution to one problem presented by ϵ -greedy policies. The greedy action a^* is still assigned the highest probability, but all other probabilities are ranked and weighted according to their values [1]. There are several ways of implementing the softmax action selection method, but for the purposes of this paper softmax uses a Boltzmann distribution, formalized in 7.

$$\forall a : p_t(a) = \frac{e^{Q_t(a)/\tau}}{\sum_{b=1}^n e^{Q_t(b)/\tau}} \quad (7)$$

The parameter τ in equation (2) is the *temperature* of the distribution. Low temperatures cause the values of $Q_t(a)$, and consequently those of $e^{Q_t(a)/\tau}$, to grow much larger, thus increasing their differences. This leads to assigning a high probability to the optimal action and assigning lower probabilities to suboptimal actions. Of course, terrible actions get very low probabilities, where in the limit $\tau \rightarrow 0$, softmax becomes greedy action selection [1]. . For high temperatures, the opposite is true. It is unclear whether ϵ -greedy action selection is better than softmax action selection. The performance of either method may depend on the task at hand and human factors. Experimenting with both methods will lead to more insight in both algorithms.

3 Implementation

The implementation consists of the following files:

Agents_new

This file contains implementations of the Agent class, the Prey class and the Predator class. Both the predator and the prey inherit functions of the Agent class. The Agent class contains functions any agent needs, such as a set of actions, a policy and other functions. As the predator is the agent is the agent this implementation focuses on, the predator class contains more functions than the predator class.

Helpers

This file contains many helper functions. These functions aid in computation and decision making, but cannot (and need not) be part of a specific class.

Other_objects

This file contains the Policy and Environment classes. The environment of the game as well as the rules are implemented in the Environment class. The Policy class contains the implementation of Q-Learning, Sarsa, ϵ -greedy, softmax action selection and more functions that help in determining and optimizing a policy as well as choosing an action of this policy.

Newstate

This file contains the Game class as well as a demonstration function. The Game class instantiates the game, initialized the predator and the prey and assigns policies to these. The game is run N times and the result is printed. The demonstration function also performs Q-Learning, Sarsa and Monte Carlo. It also uses ϵ -greedy and softmax action selection. The results are printed in the command line and graphs are used for analysis.

README

Contains instructions on how to run the program.

4 Analysis

This section discusses the results of the implementations. In order to display and compare results, graphs are used. The title describes which parameters are analysed and the legend shows which color represents which setting of said parameters.

4.1 Q-Learning

4.1.1 Different discount factors

First, the effects of different discount factors on Q-learning is analysed. The results are displayed in the figure below.




Figure 5: Q-learning using different discount factors

As shown, the higher the discount factor, the better the result. The height of the discount factor determines the importance of future rewards. A low discount rate tries to achieve a high immediate reward and a high discount rate tries to achieve a high overall reward. Since reaching each state yields a reward of 0, except for the goal state which yields a reward of 10, the future reward must be maximized. Therefore, this behaviour is expected.

4.1.2 Different learning rates

It is possible to set different learning rates in Q-learning. The learning rate determines to which extent the new information will overwrite the old information. The following graph shows how important the relationship of old and new information is to achieve an optimal policy.




Figure 6: Q-learning using different learning rates

The figure above shows that for each learning rate between, and including, 0.1 and 0.5 leads to conversion. However, a low conversion rate leads to slower conversion than a high conversion rate. It is interesting to note that the learning rates of both 0.4 and 0.5 appear to converge at about the same speed. Remember that the learning rate determines to which extent the new information will overwrite the old information. It appears that, even when storing little new data, this leads to convergence. The more new data overwrites the old data the quicker the algorithm converges. Intuitively, this seems to be correct. However, overwriting all old data might lead to less than optimal behaviour. As stated before, learning rates of 0.4 and 0.5 seem to yield the same result. It is well possible that this is the limit until which the learning rate factor is beneficial to the algorithm. As this lies beyond the scope of this report, this can be researched in the future.

4.1.3 The effect of ϵ

As Q-learning is a learning algorithm, exploration and exploitation must be balanced well to find the optimal behaviour. The following graph shows different settings for ϵ .




Figure 7: Q-learning using different epsilon values

As shown in the figure above, the lower the ϵ value is, the quicker the algorithm finds the optimal path. This can be expected since, as stated in the theory, an ϵ value of 0 leads to a greedy algorithm. What is interesting is that an ϵ value of 0.9 performs quite well compared to a completely exploratory policy, which seems to lead to random results. However, the value of 0.9 still performs a lot worse than more exploitative policies, which is expected because the environment is static (with only the prey exhibiting *slightly* stochastic behavior) and thus after a while, exploration stops being useful.

4.1.4 Softmax action selection

Besides ϵ -greedy action selection, softmax action selection was implemented. The graph and analysis can be found below.




Figure 8: Q-learning using different temperatures in Soft-max

The graph shows that all results converge. It shows that the lower the temperature, the better the results. This is to be expected, as lower temperatures lead to a more greedy action selection, which was shown above to perform better for this problem. However, the temperature of the softmax action selection algorithm can not be zero, as that would lead the algorithm to divide by zero. This ensures both optimal action selection and exploration. When analyzing the higher temperatures, it these converge at about the same speed as the lower temperatures. However, the higher temperatures, lead to less smooth graphs. As a high temperature (e.g. τ is 1) leads to an equiprobable action selection, the algorithm both explores and exploits to a certain level (going towards greedy action selection). This causes the algorithm to converge, but slower (compared to lower temperatures such as τ is 0.1) to catch the prey in an episode.

4.1.5 ϵ -greedy vs. softmax

Now let's look at the results between softmax and ϵ -greedy. As stated in the theory, the difference in effect between softmax action selection and ϵ -greedy action selection is unknown. The performance of either relies on the goal of the implementation, as well as human factors. Therefore, it is imperative to research the effects on this implementation. The ϵ value chosen for the ϵ -greedy algorithm is the same as the temperature, τ . This means that both algorithms behave greedy, but still explore from time to time.




Figure 9: Q-learning ϵ -greedy and softmax action selection

The figure shows that both softmax and ϵ -greedy action selection converge at the same speed. However, softmax action selection eventually leads to capturing the prey in less rounds. Therefore, softmax action selection is preferred over epsilon-greedy action selection when the the prey needs to be caught in a minimum amount of timesteps.

4.1.6 Different initialization for Q-learning

It was advised to initialize the policy of the predator optimistically, when performing Q-learning. When initializing a policy optimistically, it is interesting to see what happens when a policy is not optimistically initialized. The values chosen are 15, 10, 5, 0 and -5. Also, to see the effect of learning, the ϵ -value was set to either 0 or 0.1. The figure below shows the results of ϵ value 0.1.

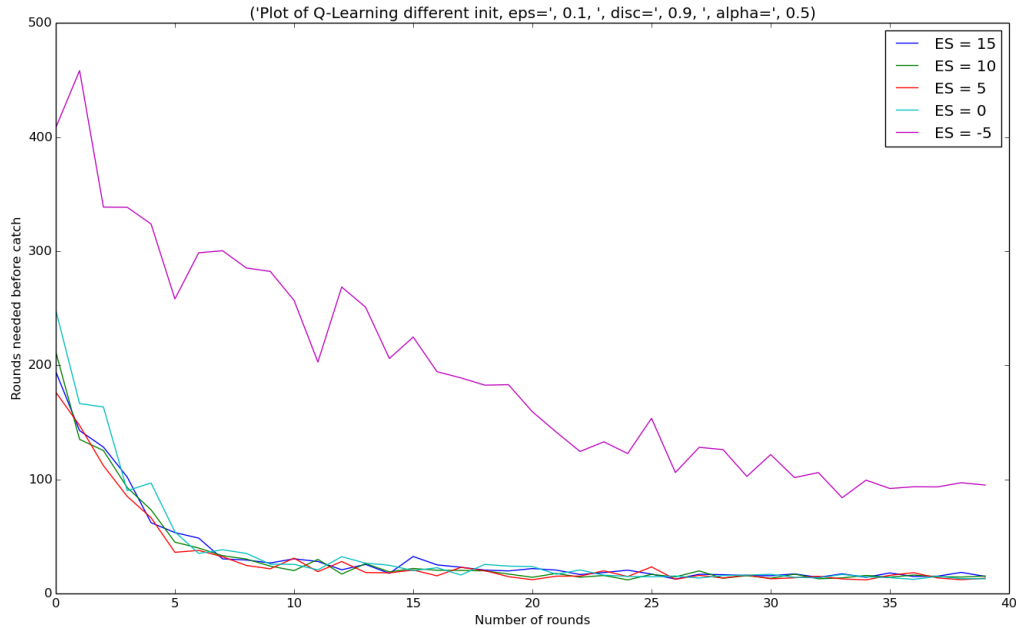


Figure 10: Q-learning with different policy initialization values and ϵ is 0.1

Each non-negative initialization appears to yield good results. Though it takes an initialization of 0 longer to converge, the policy does eventually converge to the same level as the optimistic policy. The negatively initialized policy behaves very poor. It does improve, but it takes significantly longer to converge than the non-negative initialized policies. It is expected that both the negative and zero initialization lead to slow learning. However, it is possible that the zero initialization is a border case where this is (one of) the last case(s) where the algorithm still learns enough to perform well.

Now let's analyse the same initialization values with an ϵ -value of 0.




Figure 11: Q-learning with different initialization values and ϵ is 0

This leads to similar results as before. One of the differences is that the policy initialized with 0 performs worse than before and does not converge. This behaviour can be expected. At the non-negative initializations, there are values larger than zero. Therefore, executing a greedy policy still leads to convergence. As stated, it is expected that the initialization with zero values should behave similar to the negative initialization. Again, this is not the case. This is still a border case, as stated before.

In order to determine whether the zero initialization is indeed a border case, more runs can be performed. However, this falls beyond the scope of this report. It is important to test this in the future to determine the limits of this implementation.

4.2 Learning types

Aside from Q-learning, Sarsa and on- and off-policy Monte Carlo were implemented. It is interesting to see the difference in performance between the algorithms. The results of the first three are shown in the figure below.




Figure 12: Different learning types set out against one another

The figure shows that all algorithms converge. Q-learning converges quickest, with Sarsa not much slower. On-policy Monte Carlo also converges quickly, but it starts with the worst Q-values. It also shows that of all algorithms, on-policy Monte Carlo converges to the least optimal number of rounds that lead to convergence.

Q-learning is an off-policy learning algorithm that doesn't care about the reward when learning. Sarsa does, as it learns on-policy. This also means that Sara needs to be more careful when learning, factoring in the possibility of receiving a negative reward. Therefore, Sarsa will explore less than Q-learning, leading to Q-learning learning more than Sarsa.

The following table shows how many rounds were needed, per learning type on average, for the prey to catch the predator. This was calculated for the first 1000 runs and the last 1000 runs. This means that the learning in the very beginning and the learning in the very end are set out against each other.

	Avg nr/rounds (first 1000)	Avg nr/rounds (last 1000)	STD (first 1000)	STD (last 1000)
Q-learning	88.68	14.52	124.39	18.84
Sarsa	80.61	14.70	109.74	22.45
ONMC	79.68	15.78	112.11	21.64

Table 1: Average # rounds and standard deviation of learning algorithms

Here it is concluded that Q-learning eventually leads to needing the least amount of time steps for the predator to catch the prey, although it starts worse. Sarsa performs similarly, though it starts exploiting earlier, so the average number of rounds is already lower in the beginning. In the end, the number of time steps any algorithm needs lie close to one another. This shows that all these algorithms are quick to converge and yield similar results. Therefore, either algorithm can be used if there is no reward given for reaching the goal state in a minimal amount of time steps. Also, Q-learning yields less deviation after convergence. This shows that Q-learning performs best out of these algorithms.

A late addition to the implemented learning algorithms is the off-policy Monte Carlo learning algorithm. It is therefore shown in a separate graph in Fig. 13. As you can see in comparison to the other learning algorithms in Fig. 12, the off-policy Monte Carlo starts at an equivalent amount of rounds before the prey is caught as the on-policy Monte Carlo method. The graph also shows that the off-policy Monte Carlo takes more time to converge than the other algorithms: convergence seems to take place at approximately 2000 rounds for off-policy Monte Carlo and for the other learning algorithms at approximately 500 rounds. This result is expected this algorithm learns only from the latest states of episodes, after the last non-greedy action. In this game, it is rare to have all subsequent actions greedy and this results in slower learning. Therefore, the convergence of the off-policy Monte Carlo is slow. This makes the learning of games that include long episodes inefficient.

The point of convergence for the number of round until the prey is caught for off-policy Monte Carlo also appears to stay at a slightly higher value than the other algorithms. This is possible due to the late convergence as the number of rounds needed to catch the prey may still decrease slightly after 4000 runs and will converge to the same values as the other learning algorithms. Therefore, it is recommended to run the algorithm with a much larger amount of episodes in future work.

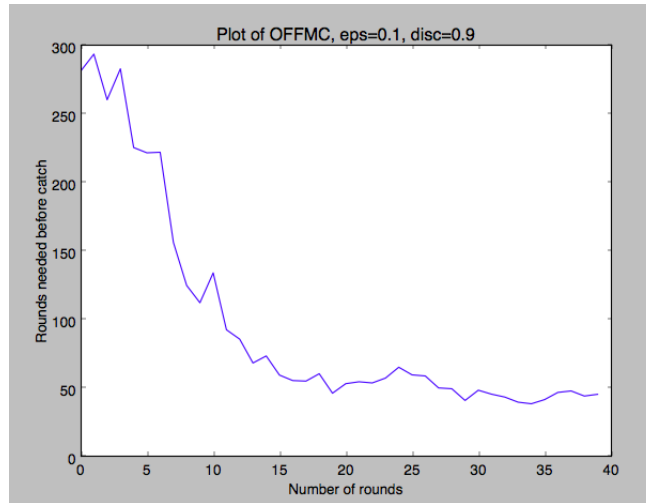


Figure 13: Off-policy Monte Carlo for 4000 runs

5 Conclusion

Q-learning is an effective method to learn the Q-values of the states of environments. This leads to learning about the states of an environment and planning to take the optimal action. Several parameters need to be tuned to achieve optimal performance. These parameters are the learning rate (α), the discount factor (*gamma*), the action selection method, ϵ , τ and the behavioural initialization. Neither Sarsa, nor on- or off-policy Monte Carlo methods have performed better for the problem at hand, which is as expected. However, if there is a high negative reward for certain cases, Sarsa should perform better.

5.1 Parameter settings

The effect of the discount factor on Q-learning shows that a high discount factor (0.9 in this case) leads to the best results. From the experiments, it is expected that a learning rate that balances old and new information (either 0.4 or 0.5 for this problem) is best.

For the problem at hand, relatively greedy methods worked quite well, which is expected for a static environment. Softmax and ϵ -greedy converged at around the same speed, but softmax performed a bit better after convergence, due to its ranking of suboptimal actions.

When using optimistic initialization, the case of $\epsilon = 0.1$ converged for all initial Q-values $\in \{-5, 0, 5, 10, 15\}$. This is expected, because exploration happens regardless due to ϵ . However, the case of $\epsilon = 0.0$ only converged for values ≤ 0 . It seems that for this problem, 0 is a tipping-point with regards to optimistic initialization.

6 Future work

In the report some results lead to more questions or needed more research than was conducted. Some of these questions were already mentioned in the report and are repeated here.

6.1 Encoded version

Currently, the implementation is not encoded. This leads to a long runtime and bad convergence for off-policy Monte Carlo. By implementing state space encoding, the runtime would decrease significantly. In state space encoding, the predator will try to minimize the distance between itself and the prey. In order to do so, the predator would calculate the distance to the prey and take action accordingly to minimize the distance to the prey. This leads to a reduced state space, leaving less states to visit and analyse.

6.2 Learning rate

The learning rate was researched between (and including) the values 0.1 - 0.5. As shown in the analysis, the learning rates of 0.4 and 0.5 seem to have similar effect. This leads to the theory that this is a turning point and rewriting more old information will actually have a negative effect on the agents' learning. Further research will show what actually happens and an optimal learning rate can be decided upon.

6.3 Varying grid initialization

The varying grid initialization showed that the grid initialization lower than 0 leads to unexpected behaviour when using an explicitly greedy policy. This appears to be the a turning point. In other words; any lower grid initialization value will probably lead to non-convergence. Also, it is interesting to find out if grid initialization with 0 is indeed a turning point or if this sometimes leads to non-convergence as well. Further research will lead to answers to these questions.

6.4 Multiple experiments

In hindsight, performing multiple experiments with the same settings would give more insight in the performance of the implementation. Thus far, all tests were run once and conclusions were drawn. More certainty about the implementation and performance of the algorithms can be achieved by averaging all results over multiple runs.

6.5 Exploration-Exploitation optimization

In order to optimize the exploration-exploitation trade-off, different techniques can be implemented. Upper confidence methods use a upper confidence bound on the probability that an action is optimal which can be used to limit the regret for a policy [1]. Another possible implementation are PAC-methods. For example, the V-max [4] algorithm uses optimistic view on state-action pairs experienced viewer than m times and adjusts this based on experience to perform efficient exploration and exploitation. Bayesian Q-learning [3] can also be used to optimize the exploration-exploitation trade-off as it keeps track of probabilities over Q values and chooses exploration and exploitation actions on basis of the expected value of the information retrieved from taking this action. These algorithms can possibly improve the agents' learning and would be interesting to implement.

7 Files attached

- newstate.py
- agents_new.py
- other_objects.py
- helpers.py ...
- README

8 Sources

References

- [1] Peter Auer, Thomas Jaksch, and Ronald Ortner. Near-optimal regret bounds for reinforcement learning. In *Advances in neural information processing systems*, pages 89–96, 2009.
- [2] Andrew G Barto. *Reinforcement learning: An introduction*. MIT press, 1998.
- [3] Richard Dearden, Nir Friedman, and Stuart Russell. Bayesian q-learning. In *AAAI/IAAI*, pages 761–768, 1998.
- [4] Karun Rao and Shimon Whiteson. V-MAX: Tempered optimism for better PAC reinforcement learning. In *AAMAS 2012: Proceedings of the Eleventh International Joint Conference on Autonomous Agents and Multi-Agent Systems*, pages 375–382, June 2012.